



# Vectorization-Based GPU Optimizations for Walberla

September 21, 2022 | Shuxin Zheng | Jülich Supercomputing Centre

# Part I: Introduction

# What is Walberla?

## Welberla

Is a modern open-source software **framework** that supports complex **multiphysics simulations**, and that is specifically designed to address the performance challenge in CSE: exploiting the full power of the largest supercomputers for a wide class of scientific research questions. WelBerla's main focus are computational fluid dynamics simulations with the **lattice Boltzmann method (LBM)**. It therefore offers a wide range of state-of-the-art LBM models, together with a variety of utility and usability functionality.

# Why is LBM important?

## Parallelizability

The LBM compute kernels make use of various node-level optimization techniques. By coupling the LBM with the RPD module (high efficient rigid particle dynamics), multiphysics simulations can be realized on a **large scale and without compromising on the performance or scalability**. In a word, its important for **parallelizability**.

# How does WalBerla achieve peak performance?

## Code generation

HPC software usually has to be modified extensively in order to make full use of new hardware architectures. WalBerla employs **code generation techniques** to generate time-critical numerical kernels from a high-level, domain-specific formulation. WalBerla used the *pystencils* metaprogramming project to generate highly efficient stencil codes for CPUs and GPUs based on a common high-level, symbolic description. As well as *lbmpy*, a code generation package that supports a wide variety of different lattice Boltzmann methods. Backends then generate C/C++, CUDA, OpeNCL code or LLVM IR from this representation. Kernel code generation has already been applied successfully to WalBerla based LBM simulations and phase-field simulations of alloy solidification on CPUs and GPUs.

## Part II: The S2A Test Case

# Pre-knowledge

## Block Forest

It is responsible for the domain partitioning and acts as a container for all data needed by the simulation. It manages the domain partitioning into blocks, independent of data structures required for the actual simulation. A figure to illustrate is shown below:



# Test Case

## S2A Test Case

The S2A test case first reads the geometry of a car from a .obj file. From this obj file, the size of the domain is determined. The size of the domain is then adapted to fit the number of cells per block. This means it is extended in all directions such that a natural number of blocks consisting of the given number of cells per block would fit for the direction. In this way, the number of blocks for the block forest is calculated and distributed to the processes. Technically, there is no need to compute the block inside the domain of the car.



# A general idea

## FlagField

However, it is not very easy, especially in Walberla, to ignore the block inside the domain of the car, due to the way the mesh is implemented. Before developing the FlagField, we need to perform all computations both inside and outside the domain of car. By implement the FlagField, we hope to reduce the cose of computing kernels by not computing the domain inside the car.

# the MLUPS table

## Master without profile

739.293 MLUPS

(million lattice cell updates per second)

739.293 MLUPS / process

734.668 MFLUPS

(million fluid lattice cell updates per second)

734.668 MFLUPS

11280.7 time steps / second

## Master with profile

633.677 MLUPS

(million lattice cell updates per second)

633.677 MLUPS / process

629.713 MFLUPS

(million fluid lattice cell updates per second)

629.713 MFLUPS / process

9669.15 time steps / second

## FlagField without profile

890.257 MLUPS

(million lattice cell updates per second)

890.257 MLUPS / process

884.688 MFLUPS

(million fluid lattice cell updates per second)

884.688 MFLUPS

13584.2 time steps /second

## FlagField with profile

707.113 MLUPS

(million lattice cell updates per second)

707.133 MLUPS /process

702.689 MFLUPS

(million fluid lattice cell updates per second)

702.689 MFLUPS / process

10789.7 time steps / second

# SIMD instructions

## CUDA MATH API

By developing the FlagField, we reduced the cost of computing kernels by not computing the domain inside the car. However, putting an 'if' condition in a GPU-based for loop is not optimal due to parallel nature of GPU threads. The motivation for using SIMD instructions here is to reduce the penalty of that 'if' condition.

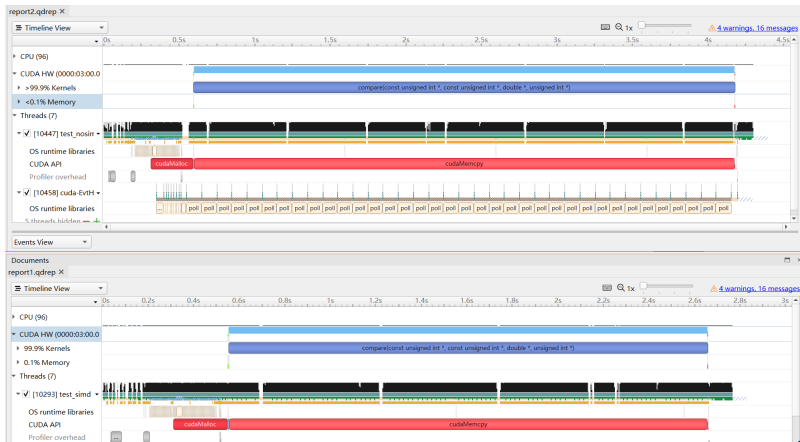
# The toy test

By developing the SIMD, we modify the origin comparing in 'if' to use the function '`__vsetne4()`' in CUDA MATH api. The function performs per-byte (un)signed comparison. It splits 4 bytes of each argument into 4 parts, each consisting of 1 byte, and returns 1 if  $a \neq b$ , else return 0.

So to be familiar with this function, and also verify the performance, I developed two version of the testing case: First fill three arrays with random unsigned int value, then call a kernel function, which do some calculation(for enlarge the GPU time) under the 'if' condition(using the *Logistic Map* here, especially  $4 * x * (x - 1)$ , where  $x$  is the random filled unsigned int value). The 'if' command comparing the two random filled array. By develop two version of tests, we get the results:

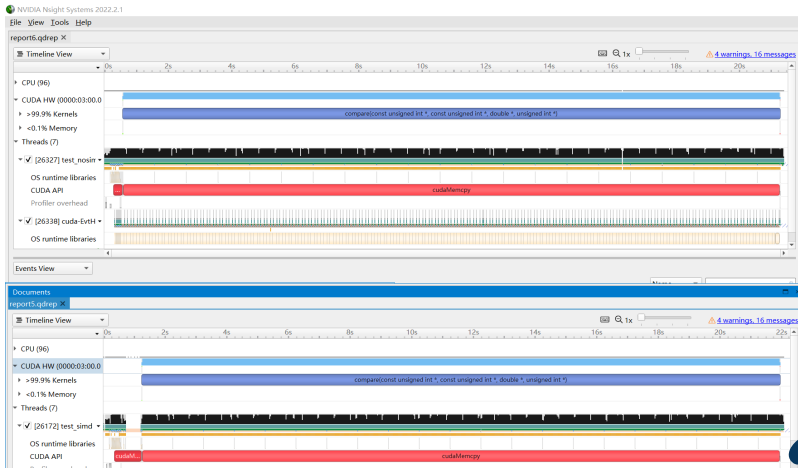
# Results1

By put the data array with size  $1024 \times 1024$ , a for loop with iteration time  $1e5$ , I got about 1.5x speed up when put the 'if' condition inside the loop.



# Results2

By put the data array with size 1024\*1024, a for loop with iteration time 1e6(for ensure and emulate the performance fluctuations), no obvious difference shown in case I put the 'if' condition outside the loop



# Current test

With the assumption that the kernel function in WalBerla is called large more than once, I adapting this SIMD modification into the origin branch. The preliminary tested performance shows a 1.11x speed up:

## Master with profile

633.677 MLUPS

(million lattice cell updates per second)

633.677 MLUPS / process

629.713 MFLUPS

(million fluid lattice cell updates per second)

629.713 MFLUPS

9669.15 time steps / second

## SIMD with profile

704.314 MLUPS

(million lattice cell updates per second)

704.314 MLUPS / process

699.908 MFLUPS

(million fluid lattice cell updates per second)

699.908 MFLUPS

10747 time steps /second

# Future Work

## work need to be finished

- Improve the test for including SIMD command on basic code
- Deploy the SIMD into the FlagField version
- Test the performance for SIMD inside the FlagField





# Vectorization-Based GPU Optimizations for Walberla

September 21, 2022 | Shuxin Zheng | Jülich Supercomputing Centre